

Implementation of Concurrent Hash Tables in Java

Shawn Guydeene

Dalton Kajander

Blake Celestian

Abstract—Concurrent programming is using multiple processes at the same time, often reducing total time waiting for operations to complete. A concurrent hash table is a hash table that utilizes concurrent programming in order to reduce the time it takes to do multiple operations, such as the `add()`, `search()`, and `remove()` methods hash tables utilize. It does this by doing different operations at the same time instead of one after another, waiting for the previous to finish before the next can even start. This report is about attempting to create our own concurrent hash table in the Java language instead of using Java's own `HashTable` class. We will be achieving this by using function keywords like `synchronized` and manually locking and unlocking, experimenting with both options. We will show evidence of reduced time by running the same commands on an unthreaded custom `HashTable` and on a custom concurrent `HashTable`.

Index Terms—Concurrent Programming, Hash-table representations

I. INTRODUCTION

The problem that our group chose to tackle, is how one would parallelize a hash table, and see how much of a performance increase, if any would arise from such an implementation. This paper seeks to explain how hash tables function as a data structure, previous attempts at parallelizing them, our group's attempt at creating a concurrent hash table, and experimental data collected showing the effects of parallelizing on runtime for the program.

II. RESEARCH

A. What are Hash Tables?

Hash tables are a data structure that utilizes an associative array, which is a structure that associates keys to specific values, in order to store data. The way in which a hash table uses these keys, is by entering them into a hash function. The hash function will then compute the array index in which the value associated with the key will be stored. The hash function is naturally the most important portion of the hash table's construction. The most commonly used hashing method uses some variation of modular hashing. Modular hashing works as follows: assuming there is an array of size M , where M is prime, then for any positive integer key k , compute $k \bmod M$.

Modular hashing is expandable to accepting more than just integer as keys. A common data type used as a key for hash tables would be strings. In order to accomplish this, one would simply have to treat the string as though it were a large integer. This concept of converting data types into integers is commonly used as it is much simpler than creating a custom hashing function. A well made hash function will have three

main requirements: deterministic - equivalent keys will always produce the same hash value, computationally efficient, and allows for keys to be uniformly distributed. Deterministic is only one of those requirements that is absolutely necessary for any hash function, as it what allows users to search hash tables for a specific key-value pairing [1].

There is of course, one issue with the hash function: since the hash table is not an array of infinite size, there is guaranteed to be instances where two distinct keys map to the same index when using the hash function. In order to resolve this, several techniques have been developed: linear probing, quadratic probing, double hashing, separate chaining or buckets. The basic concept for linear probing is that if the current key-value attempting to be entered into the hash table collides at position i , then attempt to place it at $i + 1$. If it collides again, keep increasing the index by 1 until it comes across an index in the hash table array that has not been used yet, ending at $i + n$ where n is the number of collisions. Quadratic probing is similar to linear probing, except instead of ending up at $i + n$, the value will be stored at index $i + n^2$, where n is the number of collisions. Double hashing, also known as rehashing, involves hashing the key a second time, with a different hash function. That result is taken, and then used as the step size for a process similar to the previous two methods, resulting in the value being placed at index $i + jn$, where n is the number of collisions, and j is the result of the second hash function. These methods, known as open addressing, work, but as the hash table begins to fill up, they result in the `add()` operation taking more time to complete due to collisions. Separate chaining, on the other hand, works by having each index of the array be a linked list. Whenever a value maps to that index, it will be added into list, which avoids the issue that open addressing presents. Buckets is a similar method to separate chaining, using an array instead of a linked list. While separate chaining and buckets fix the issues with open addressing, they present more of their own, namely that there exist the possibility that a large majority of keys will map to the same index, which leads to a large amount of time being used searching the array or list at that index for the specified key's value [2].

B. What is Concurrency?

Now given what a hash table is, this paper will look at what exactly it means for a computer program to be concurrent. Concurrency means that multiple computations will be occurring at the same time. This has become increasingly more important, with the clock speed of processors no longer increasing, it results in manufacturers adding more cores to

chips, meaning splitting up computation among those chips is the only way to make it faster. Concurrent programming has two common models: shared memory models and message passing models. This paper will only look at shared memory models as the implementation for a concurrent hash table will utilize a shared memory model. A shared memory model works by having concurrent processors read and write objects into the memory that they share [3].

III. ALGORITHMS

There are many ways to create single-threaded hash tables, so there are likewise many ways to create concurrent hash tables. If we ignore the need for resizing and the choice of hashing algorithm, we still need to decide on an approach for handling contention. We will attempt to employ methods such as course-grain locking, fine-grain locking, lazy, optimistic, and lock-free design.

IV. EXPERIMENTAL RESULTS

We begin by testing how the standard hash table will performed when run sequentially, and then when multithreaded.

V. APPENDIX

A. Challenges

One of the major challenges the group came across was learning how to do concurrent programming within Java, as most of the group member's experience lay in C++.

B. Completed Tasks

- Creating a standard hash table.
- Creating a hash table with locks.

C. Remaining Tasks

- Creating a lock-free hash table.
- More experimental testing.

REFERENCES

- [1] R. Sedgewick and K. Wayne. "Hash Tables." [princeton.edu. https://algs4.cs.princeton.edu/34hash/](https://algs4.cs.princeton.edu/34hash/) (accessed Mar. 13, 2022).
- [2] S. Mitra. "Hashing." [utexas.edu. https://www.cs.utexas.edu/mitra/csSpring2017/cs313/lectures/hash.html](https://www.cs.utexas.edu/mitra/csSpring2017/cs313/lectures/hash.html) (accessed Mar. 13, 2022).
- [3] S. Amarasinghe, A. Chlipala, S. Devadas, M. Ernst, M. Goldman, J. Guttag, D. Jackson, R. Miller, M. Rinard, and A. Solar-Lezama. "Reading 17: Concurrency" [mit.edu. https://web.mit.edu/6.005/www/fa14/classes/17-concurrency/](https://web.mit.edu/6.005/www/fa14/classes/17-concurrency/) (accessed Mar. 14, 2022).